

# SPARK EXECUTION

So far, we've interacted with Spark

1. **Using** an REPL environment

2. **Submitting** a script

These are ways for the user to  
give instructions to Spark

Spark takes the user's instructions

It translates them into tasks that  
run across a computing cluster

It returns the results to the user

How does this work?

When you run a program (using PySpark or spark-submit)

Spark prints **a bunch of messages**

First, there are a lot of messages  
initializing things

```
INFO util.Utils: Successfully started service 'sparkDriver' on port 64191.  
INFO slf4j.Slf4jLogger: Slf4jLogger started  
INFO Remoting: Starting remoting  
INFO Remoting: Remoting started; listening on addresses :[akka.tcp://sparkDriverActorSystem@  
INFO util.Utils: Successfully started service 'sparkDriverActorSystem' on port 64192.  
INFO spark.SparkEnv: Registering MapOutputTracker  
INFO spark.SparkEnv: Registering BlockManagerMaster
```



# Then you can see the **progress** of your processing instructions

```
spark.SparkContext: Starting job: saveAsTextFile at  
scheduler.DAGScheduler: Got job 0 (saveAsTextFile at  
scheduler.DAGScheduler: Submitting ResultStage 0 (MapPar  
cluster.YarnScheduler: Adding task set 0.0 with 2 tasks  
scheduler.TaskSetManager: Starting task 0.0 in stage 0.0  
scheduler.TaskSetManager: Starting task 1.0 in stage 0.0  
scheduler.TaskSetManager: Finished task 1.0 in stage 0.0  
scheduler.TaskSetManager: Finished task 0.0 in stage 0.0  
scheduler.DAGScheduler: ResultStage 0 (saveAsTextFile at  
scheduler.DAGScheduler: Job 0 finished: saveAsTextFile
```



# There seem to be **many components** involved in the processing

```
spark.SparkContext Starting job: saveAsTextFile at  
scheduler.DAGScheduler: Got job 0 (saveAsTextFile at  
scheduler.DAGScheduler: Submitting ResultStage 0 (MapPar  
cluster.YarnScheduler: Adding task set 0.0 with 2 tasks  
scheduler.TaskSetManager: Starting task 0.0 in stage 0.0  
scheduler.TaskSetManager: Starting task 1.0 in stage 0.0  
scheduler.TaskSetManager: Finished task 1.0 in stage 0.0  
scheduler.TaskSetManager: Finished task 0.0 in stage 0.0  
scheduler.DAGScheduler: ResultStage 0 (saveAsTextFile at  
scheduler.DAGScheduler: Job 0 finished: saveAsTextFile
```

What are all of these components doing?

SparkContext

DAGScheduler

YarnScheduler

TaskSetManager

SparkContext

Every Spark application  
has a driver program

DAGScheduler

YarnScheduler

TaskSetManager



SparkContext

driver program

This could be

The PySpark/Scala shell

A script

A main function  
(Java/Scala)

DAGScheduler

YarnScheduler

TaskSetManager

SparkContext

driver program

This program consists of  
**instructions** to Spark

**Load data** into RDDs

Perform **operations**  
on RDDs

DAGScheduler

YarnScheduler

TaskSetManager

driver program

SparkContext

Driver programs use a  
**SparkContext** to  
communicate with the  
Spark cluster

DAGScheduler

YarnScheduler

TaskSetManager



driver program

SparkContext

There are 2 phases involved  
in the driver program

Initial setup

Job Run

DAGScheduler

YarnScheduler

TaskSetManager

driver program

SparkContext

Initial setup

This happens when you  
launch the program (the  
shell or your script)

Job Run

DAGScheduler

YarnScheduler

TaskSetManager

# driver program

## SparkContext

### Initial setup

A Job run is initiated whenever the program has a processing task

### Job Run

Ex: An Action like  
**count() or collect()**

DAGScheduler

TaskScheduler

TaskSetManager



# Initial setup

driver program

SparkContext

Spark needs a **separate cluster manager** to manage resources across the cluster

This is a plug and play component : **Mesos, YARN** or **Spark Standalone**


DAGScheduler

YarnScheduler

TaskSetManager

# Initial setup

Mesos, YARN or Spark Standalone



When the driver program **starts**


It contacts the **cluster manager**  
through the **SparkContext**

DAGScheduler

TaskSetManager

# Initial setup

Mesos, YARN or Spark Standalone



The cluster manager in turn  
launches Java processes on  
several nodes in the cluster

DAGScheduler  
TaskSetManager



# Initial setup

Mesos, YARN or Spark Standalone

driver program

SparkContext

cluster manager

YarnScheduler

## Executors

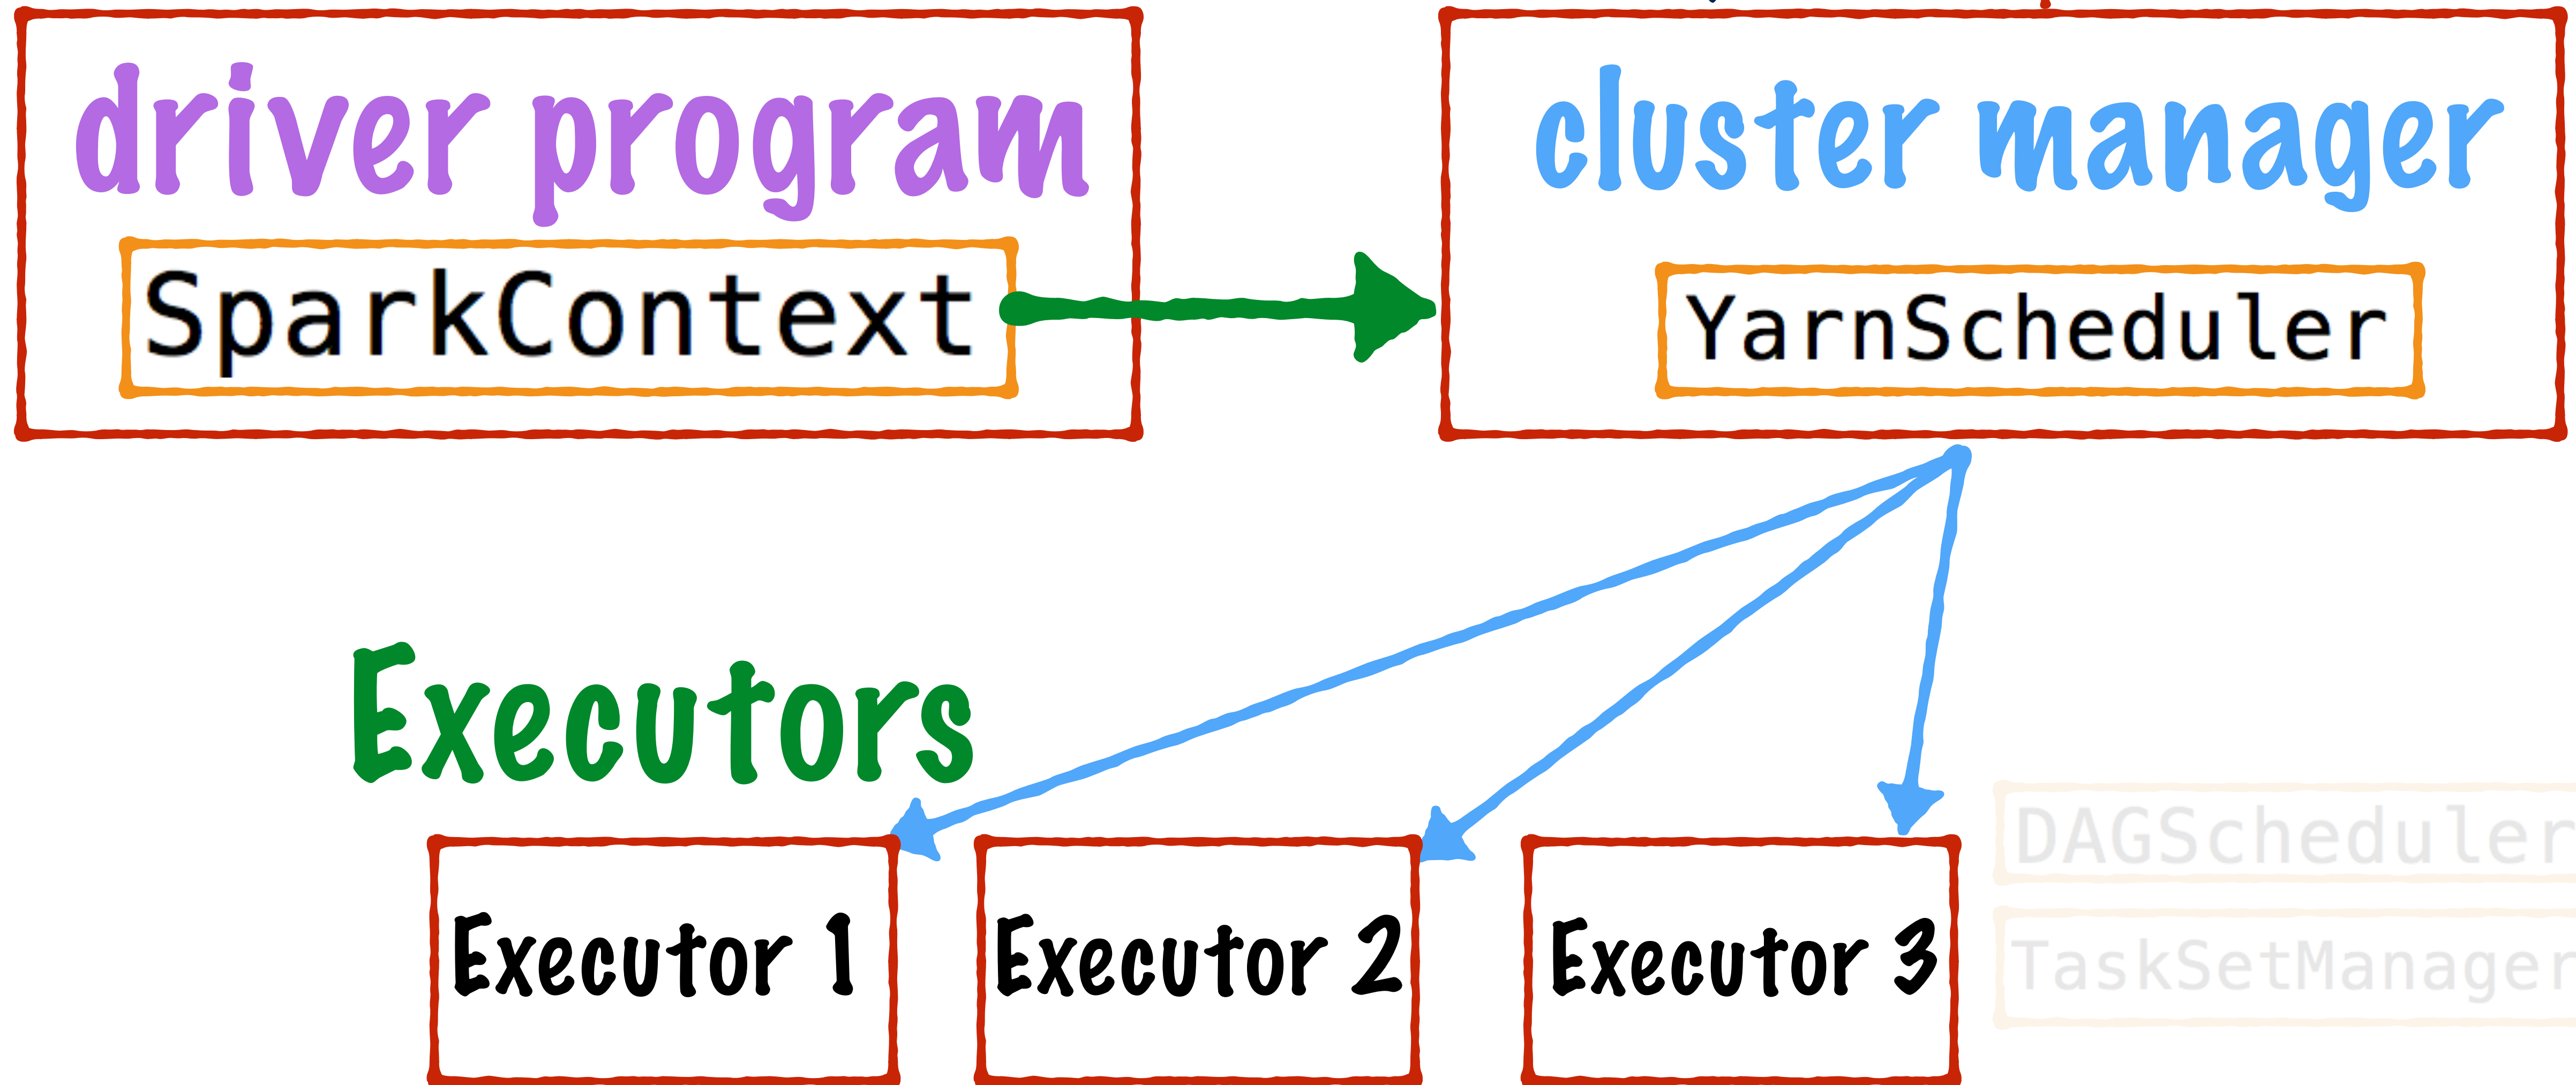
The cluster manager in turn  
launches Java processes on  
several nodes in the cluster

DAGScheduler

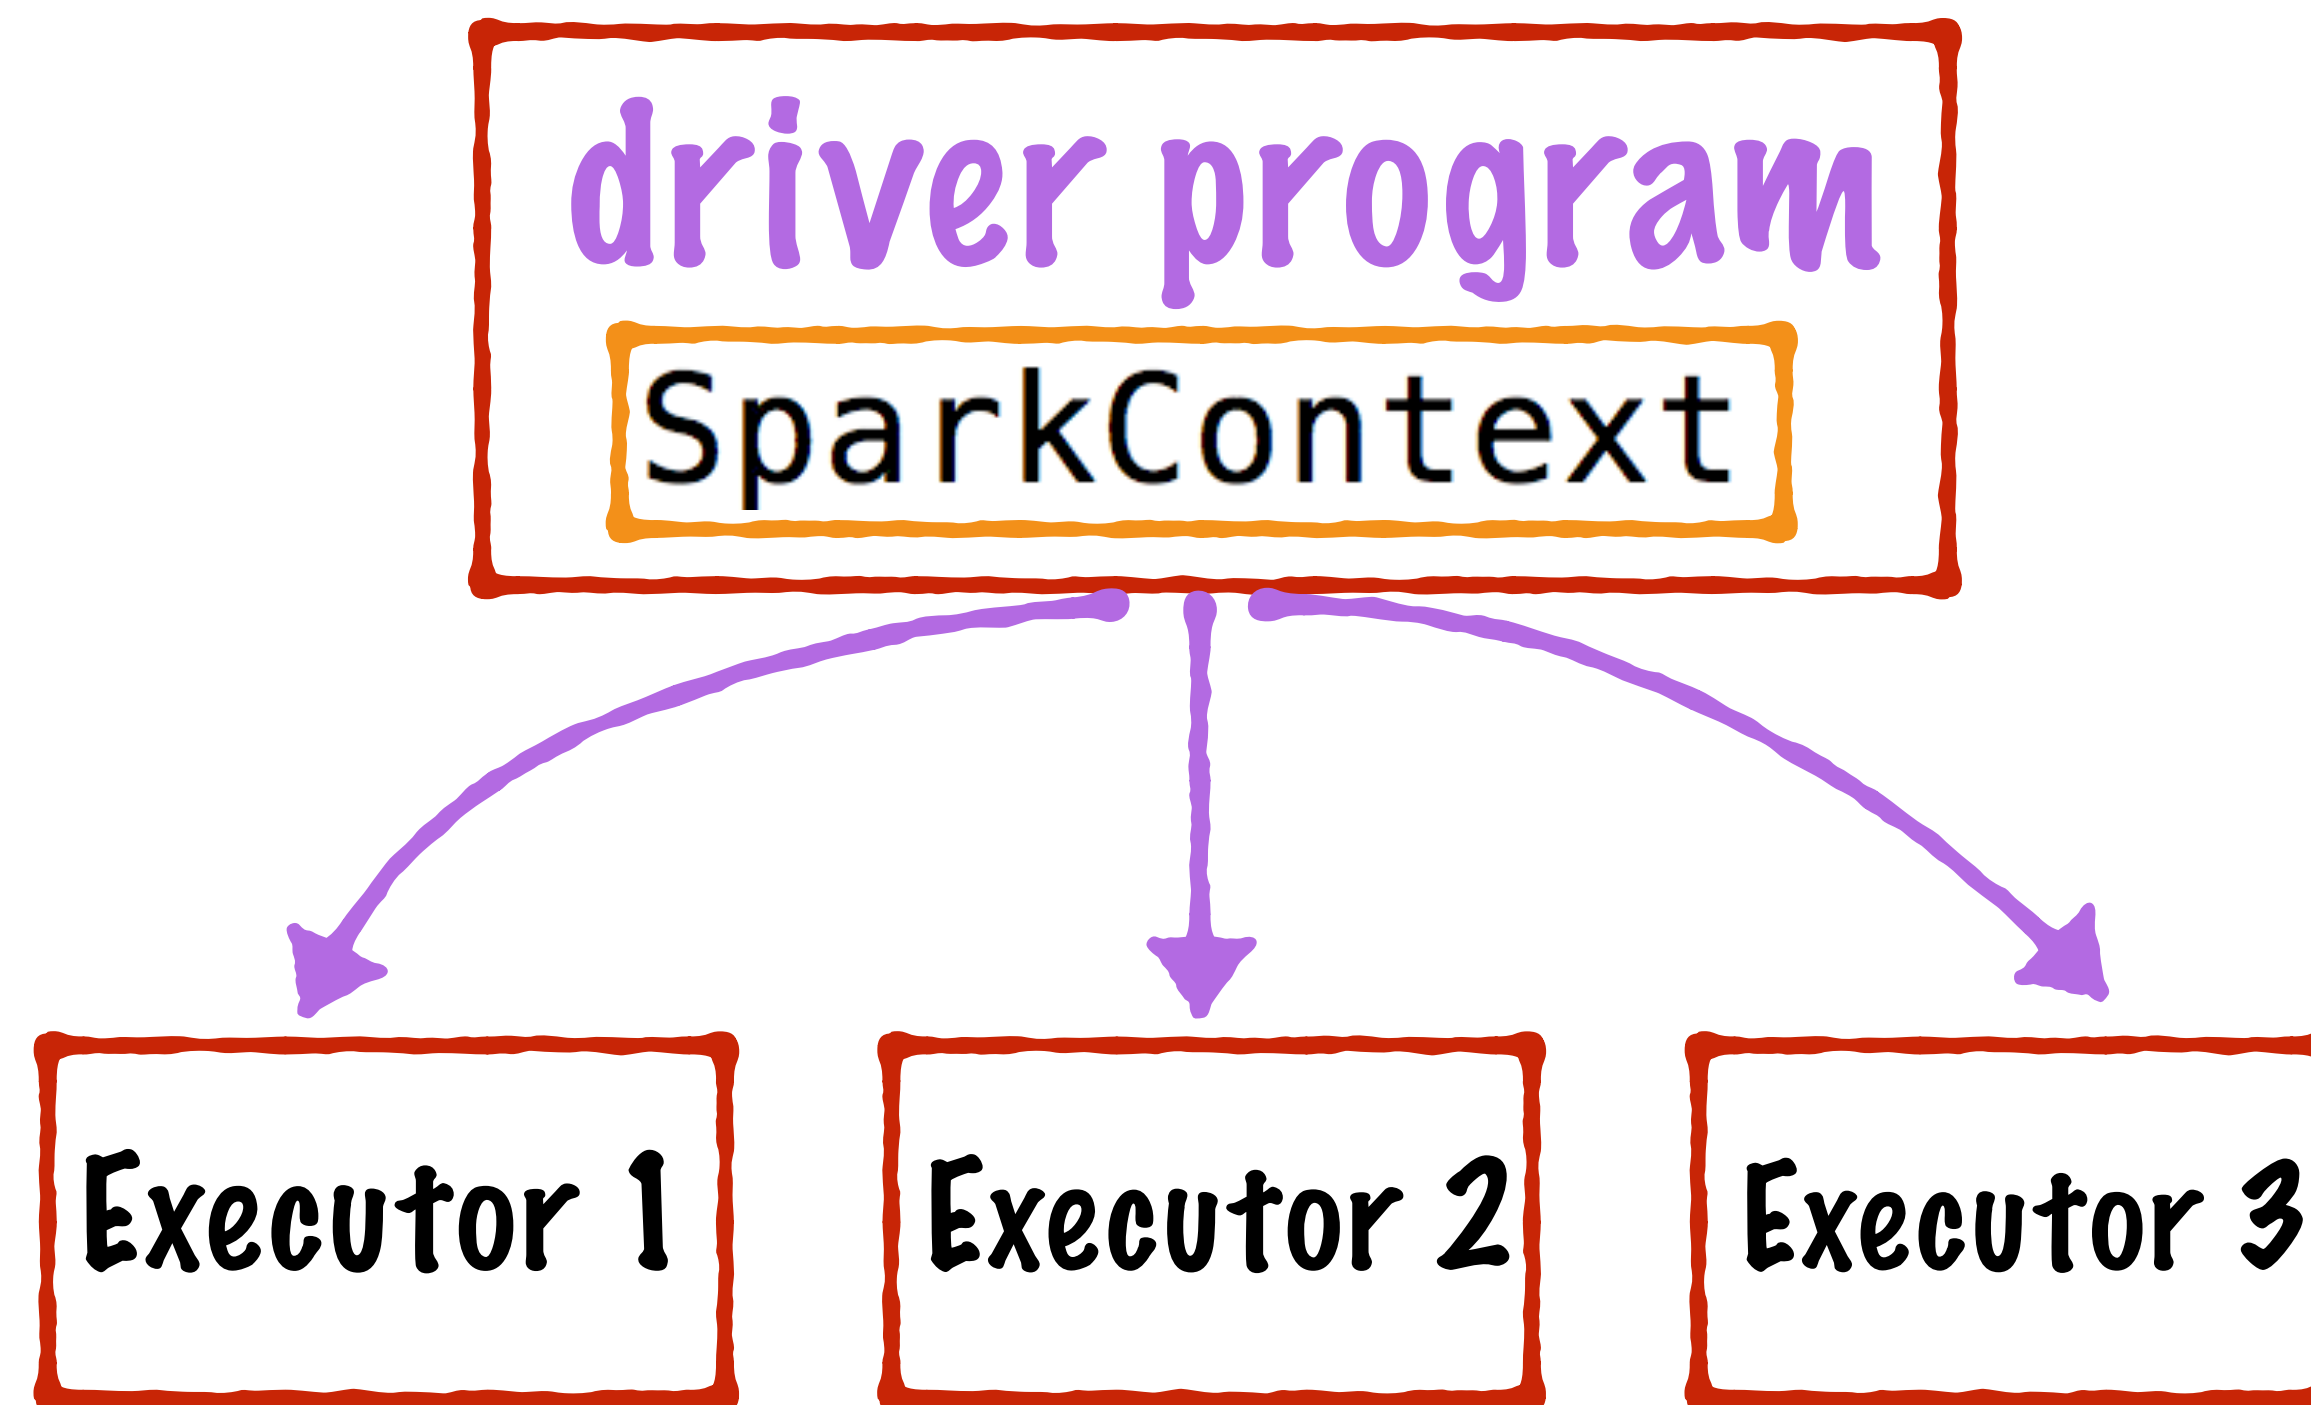
TaskSetManager

# Initial setup

Mesos, YARN or Spark Standalone



# Initial setup



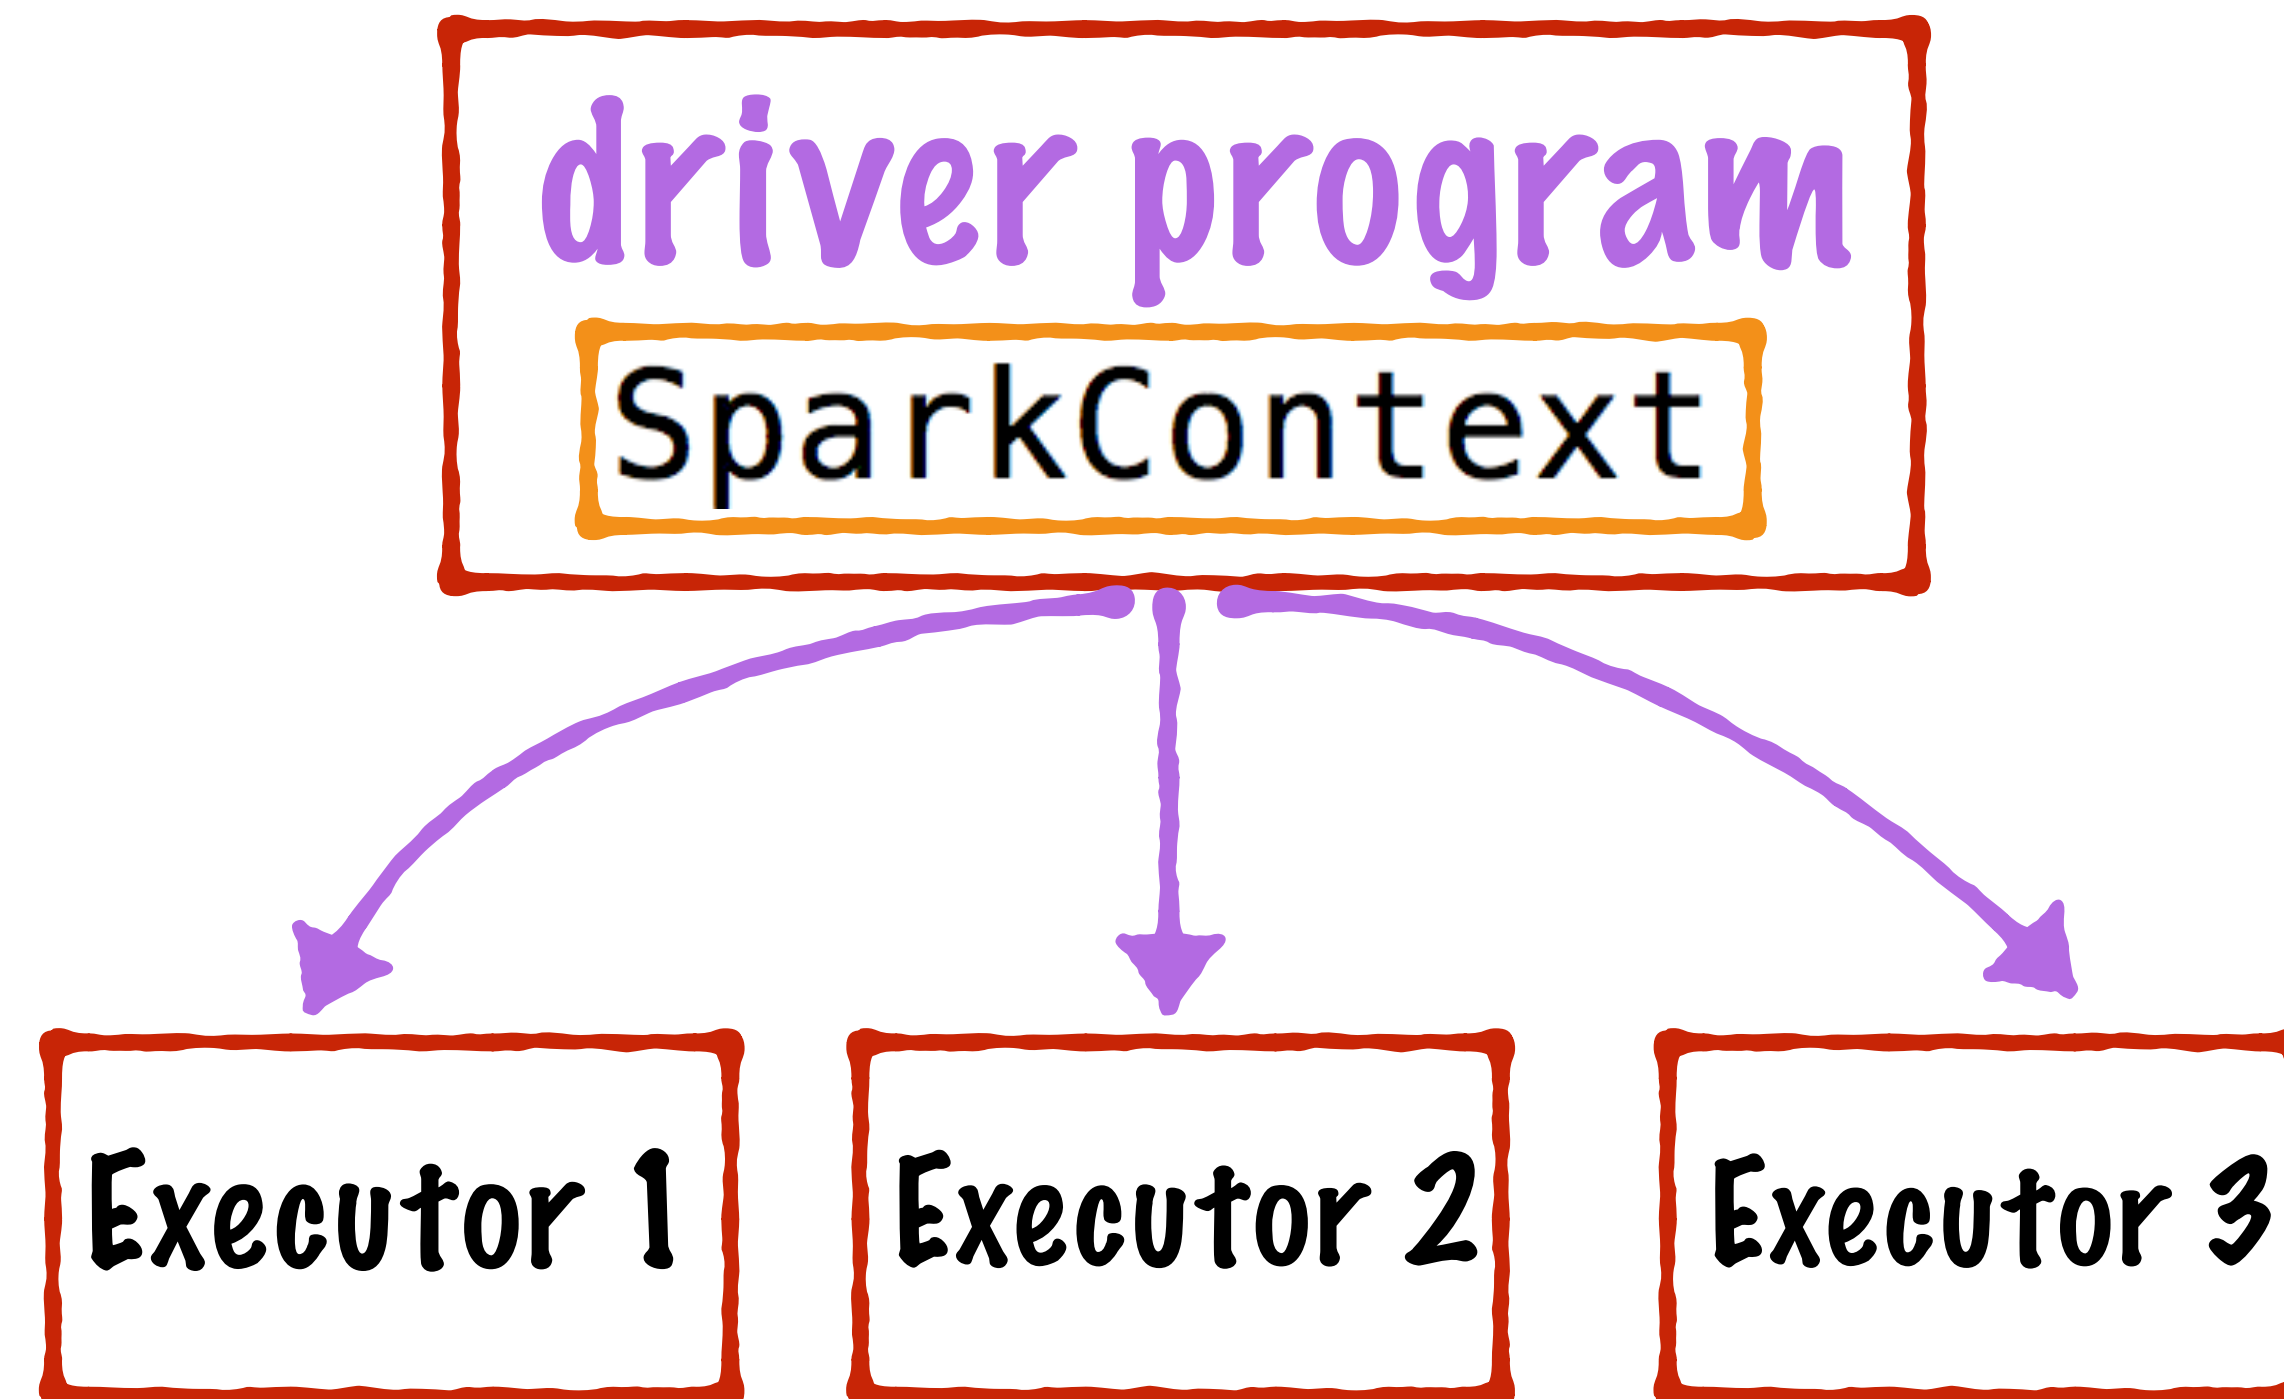
Once the executors are launched  
they **register themselves** with  
the driver program

DAGScheduler

TaskSetManager



# Initial setup



The driver program is  
**now ready** for processing  
user instructions

DAGScheduler

TaskSetManager

driver program

SparkContext

Initial setup ✓

Job Run

DAGScheduler

TaskSetManager

driver program

SparkContext

# Job Run

Whenever there is an action  
(or some processing) on an RDD

A Job is initiated

When this happens,  
SparkContext passes the user  
instructions on to a Scheduler

DAGScheduler

TaskSetManager

Executor 1

Executor 2

Executor 3


# Job Run

driver program


SparkContext

DAGScheduler

Job



Stages



Tasks

The scheduler breaks  
down the **Job** into  
**smaller units of work**

TaskSetManager



# Job Run

driver program

SparkContext

DAGScheduler

Job

```
flightsParsed.map(lambda x:x.distance).reduce(lambda x,y:x+y)
```

Stages

Stage 0

```
map(lambda x:x.distance)
```

Stage 1

```
.reduce(lambda x,y:x+y)
```

Tasks

Tasks on individual partitions of the RDD

TaskSetManager

# Job Run

The Stages and Tasks form  
a **Directed Acyclic Graph**

**DAG**

This is a standard way  
to represent a workflow

driver program

SparkContext

DAGScheduler

Job

Stages


Tasks

TaskSetManager

# Directed Acyclic Graph

In any workflow

You might  
have a bunch  
of tasks




# Directed Acyclic Graph

Some tasks are independent

They can be done in  
**parallel**

Ex: Applying a map  
operation on individual  
partitions of an RDD






# Directed Acyclic Graph


Some tasks are  
dependent on each other  
They have to be performed **serially**

Ex: In reduce,

1. individual partitions have to be processed
2. Those results are combined



# Directed Acyclic Graph



# Job Run

The Stages and Tasks form  
a **Directed Acyclic Graph**

**DAG**

This is a standard way  
to represent a workflow

driver program

SparkContext

DAGScheduler

Job

Stages

Tasks

TaskSetManager

# Job Run

driver program

SparkContext

DAGScheduler

Tasks

Tasks are the smallest unit of work in a Spark Job

Each task is **assigned to an Executor** by the Task scheduler

Executor 1

Executor 2

Executor 3

Task Scheduler



# Job Run

driver program

SparkContext

DAGScheduler

TaskSetManager

Tasks

Each task is assigned to  
an Executor by the Task  
scheduler

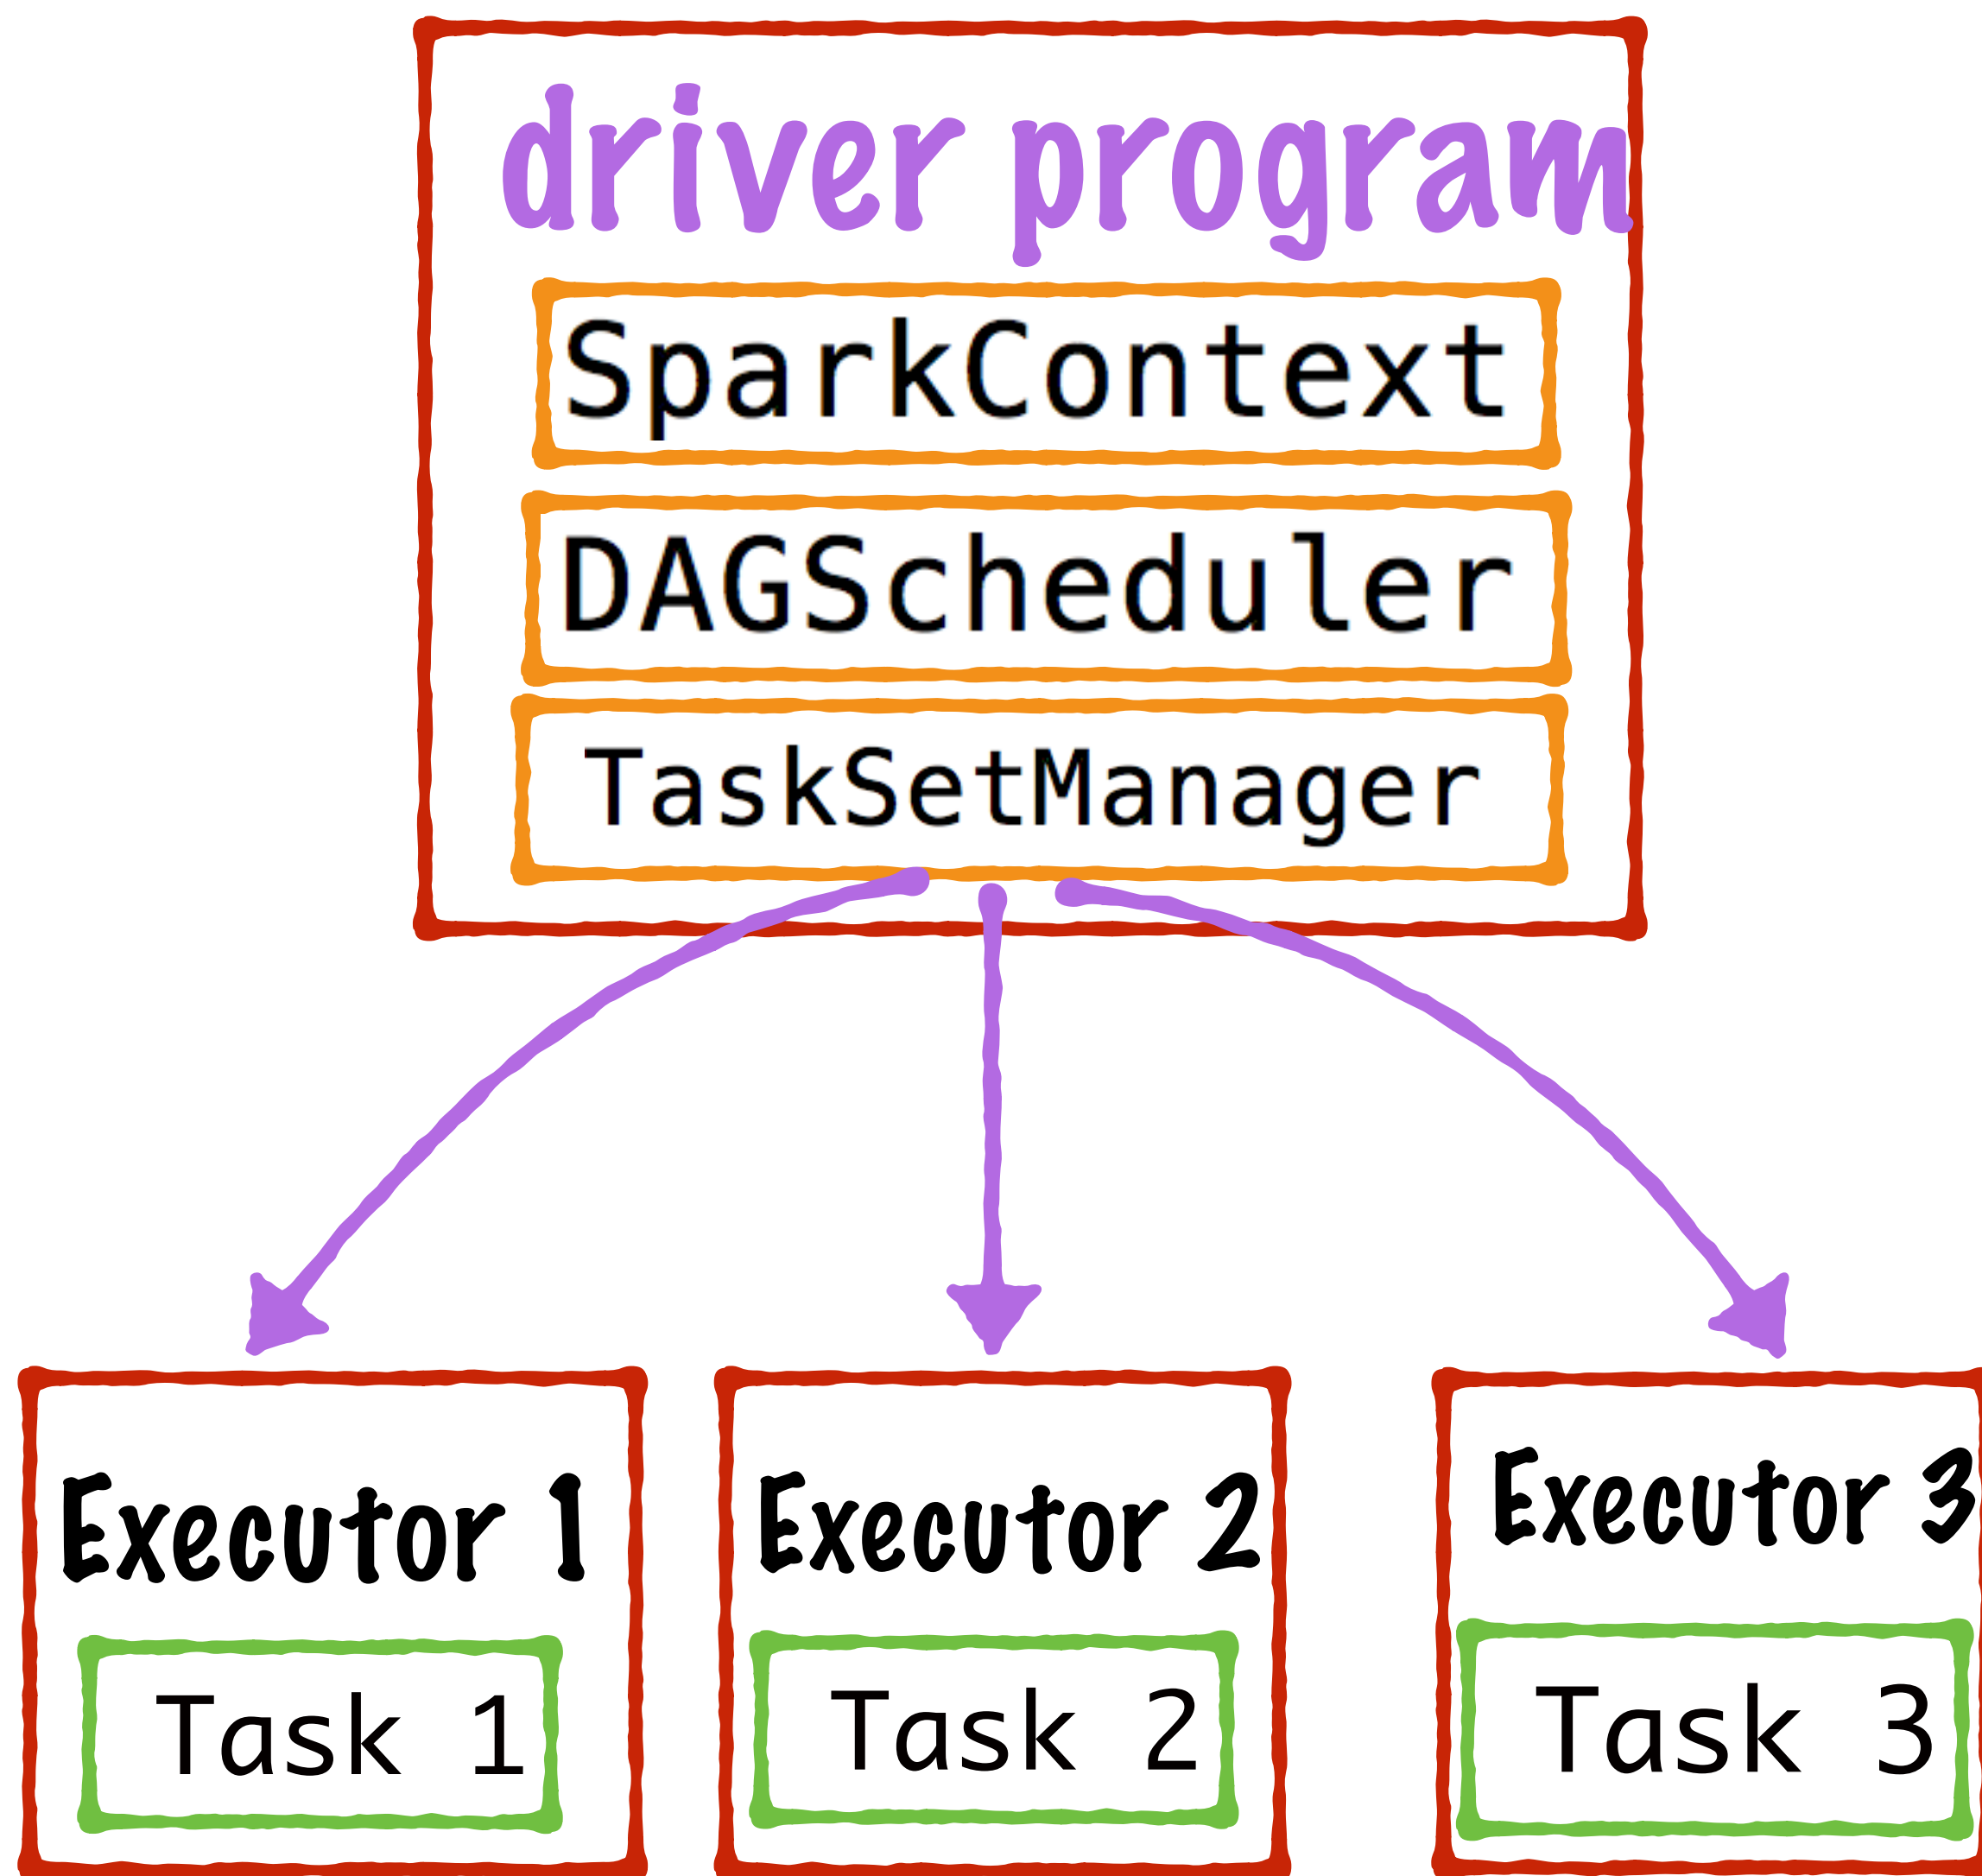
Executor 1

Executor 2

Executor 3


# Job Run

The executors run the tasks and send updates back to the driver program



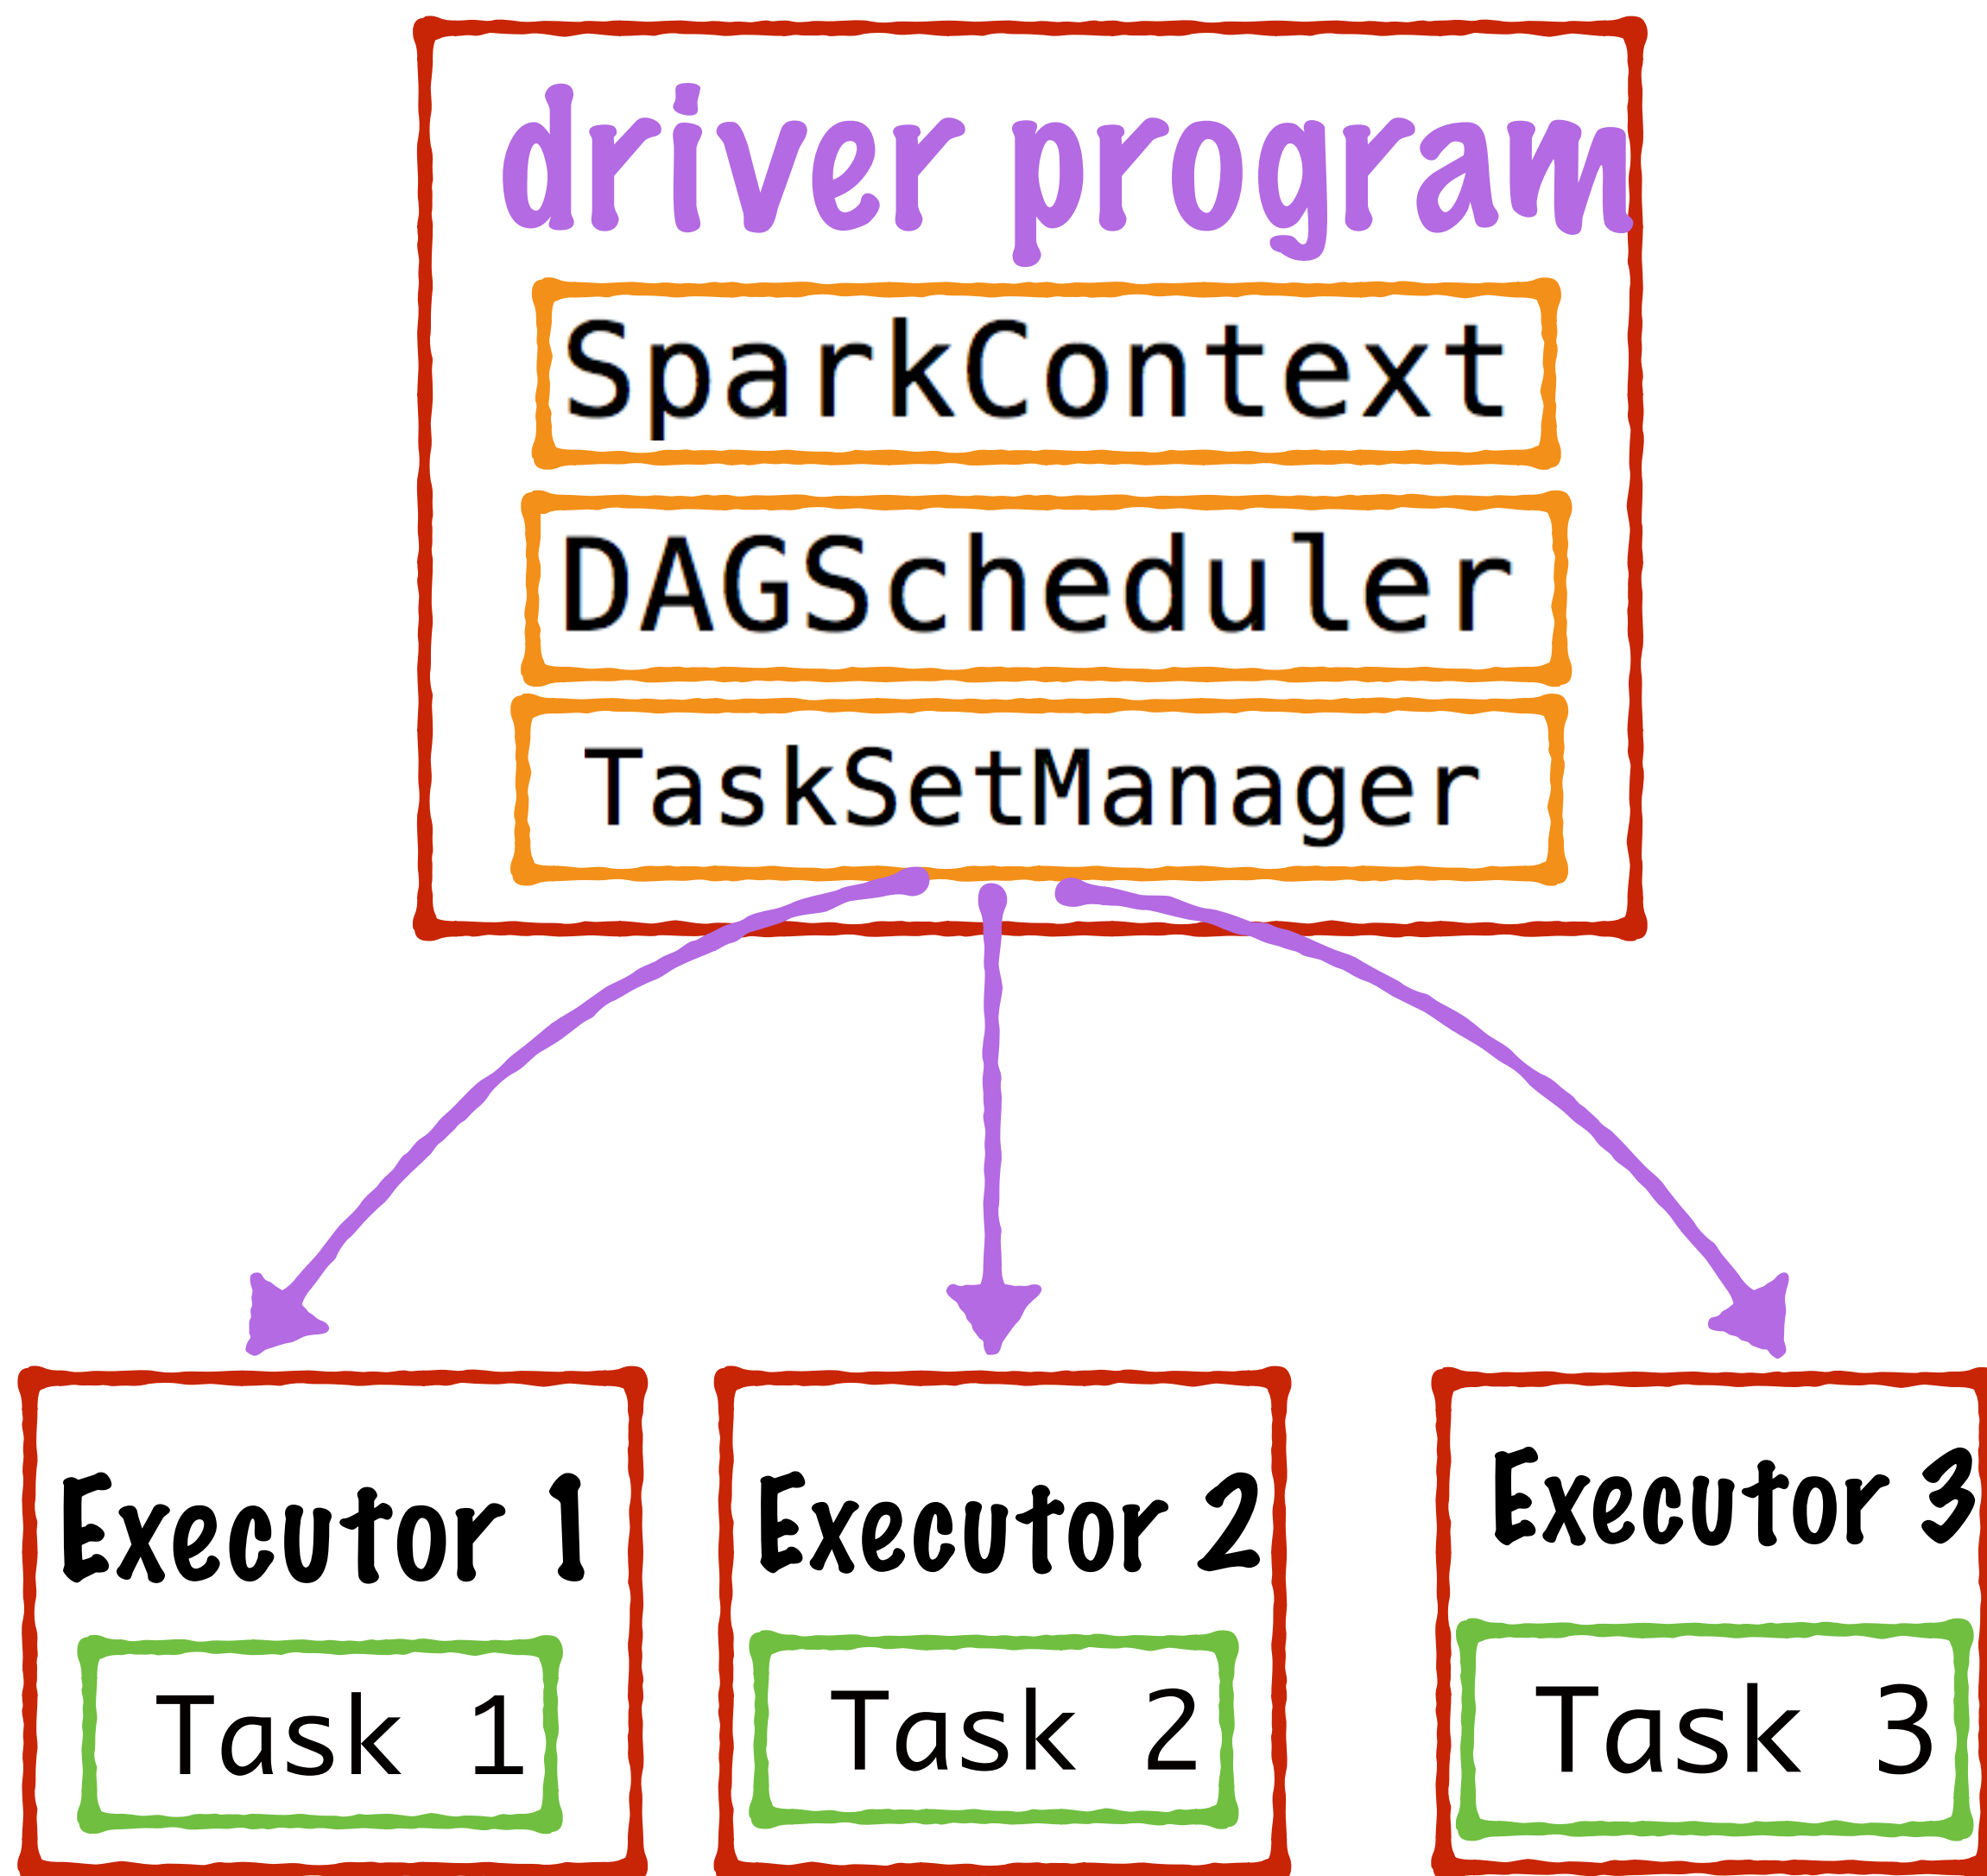
# Job Run

When all the tasks in a Stage are completed, the next Stage is launched



# Job Run


When all the Stages are completed, the Job finishes






# Job Run

Let's go back and look  
at the status messages  
Spark prints



# Job Run

Let's go back and look  
at the status messages  
Spark prints




# Job Run

driver program

SparkContext

DAGScheduler

TaskSetManager



Executor 1

Task 1

Executor 2

Task 2

Executor 3

Task 3

```
spark.SparkContext: Starting job: saveAsTextFile
scheduler.DAGScheduler: Got job 0 (saveAsTextFile)
scheduler.DAGScheduler: Submitting ResultStage 0 (
cluster.YarnScheduler: Adding task set 0.0 with 2 tasks
scheduler.TaskSetManager: Starting task 0.0 in stage 0.
scheduler.TaskSetManager: Starting task 1.0 in stage 0.
scheduler.TaskSetManager: Finished task 1.0 in stage 0.
scheduler.TaskSetManager: Finished task 0.0 in stage 0.
scheduler.DAGScheduler: ResultStage 0 (saveAsTextFile) finished
scheduler.DAGScheduler: Job 0 finished: saveAsTextFile
```



# Job Run

driver program

SparkContext

DAGScheduler

TaskSetManager

```
spark.SparkContext: Starting job: saveAsTextFile
scheduler.DAGScheduler: Got job 0 (saveAsTextFile)
scheduler.DAGScheduler: Submitting ResultStage 0 (
cluster.YarnScheduler: Adding task set 0.0 with 2 tasks
scheduler.TaskSetManager: Starting task 0.0 in stage 0.
scheduler.TaskSetManager: Starting task 1.0 in stage 0.
scheduler.TaskSetManager: Finished task 1.0 in stage 0.
scheduler.TaskSetManager: Finished task 0.0 in stage 0.
scheduler.DAGScheduler: ResultStage 0 (saveAsTextFile) is finished
scheduler.DAGScheduler: Job 0 finished: saveAsTextFile
```

SparkContext  
starts the Job

Executor 1

Task 1

Executor 2

Task 2

Executor 3

Task 3



# Job Run

driver program

SparkContext

DAGScheduler

TaskScheduler

```
spark.SparkContext: Starting job: saveAsTextFile
scheduler.DAGScheduler: Got job 0 (saveAsTextFile)
scheduler.DAGScheduler: Submitting ResultStage 0 (
cluster.YarnScheduler: Adding task set 0.0 with 2 tasks
scheduler.TaskSetManager: Starting task 0.0 in stage 0.
scheduler.TaskSetManager: Starting task 1.0 in stage 0.
scheduler.TaskSetManager: Finished task 1.0 in stage 0.
scheduler.TaskSetManager: Finished task 0.0 in stage 0.
scheduler.DAGScheduler: ResultStage 0 (saveAsTextFile) is finished
scheduler.DAGScheduler: Job 0 finished: saveAsTextFile
```

DAGScheduler takes the  
Job and constructs a  
DAG of Stages and tasks

# Job Run

driver program

SparkContext

DAGScheduler

TaskSetManager

```
spark.SparkContext: Starting job: saveAsTextFile
scheduler.DAGScheduler: Got job 0 (saveAsTextFile)
scheduler.DAGScheduler: Submitting ResultStage 0 (
cluster.YarnScheduler: Adding task set 0.0 with 2 tasks
scheduler.TaskSetManager: Starting task 0.0 in stage 0.
scheduler.TaskSetManager: Starting task 1.0 in stage 0.
scheduler.TaskSetManager: Finished task 1.0 in stage 0.
scheduler.TaskSetManager: Finished task 0.0 in stage 0.
scheduler.DAGScheduler: ResultStage 0 (saveAsTextFile) a
scheduler.DAGScheduler: Job 0 finished: saveAsTextFile
```

TaskSetManager

actually schedules and  
keeps track of the tasks